Guardrail: proving every agent action before it runs

Team [Team name] · **Team ID** [_____] · **Theme 6** AI You Can Actually Trust · **Problem 9** Catching an Agent Before It Acts on a Bad Idea

PROBLEM & THEME CHOSEN

Tool-calling agents now move money, send messages and edit records, and their failures look like success: a refund on a mistyped order number, a misquoted policy, an email to a near-miss address. These effects cannot be undone, so the check must happen before execution, against sources the model cannot talk its way around.

UNDERSTANDING OF THE CHALLENGE

The real difficulty has four parts. First, hallucinated actions are usually valid on paper: a $129.99 refund on the customer's real order passes every existence check, yet it is the tent's price when she asked about the lantern. Second, no contradiction is not proof: a blocklist of known mistakes misses the next one. Third, checking and doing must be one step: records change between check and write, and retries can execute twice. Fourth, an LLM verifier is itself an attack surface: it can be prompt-injected through the agent's arguments, or approve on evidence it invented.

PROPOSED APPROACH

A guardrail service between the agent's decision and tool execution, in a support-desk domain (refunds, emails, policy answers, stock changes): a LangGraph agent (agent → verify → await approval → execute) on Claude Opus 5.5, and a SQLite system of record. Every proposed call passes four stages:

1. A strict schema per tool rejects invented tools, parameters and malformed values.
2. Thirty deterministic checks, written as an allowlist, each return PASS, FAIL (block), UNVERIFIABLE or NEEDS_APPROVAL (both go to a person). They cover ownership of every ID; provenance (every ID, amount and quantity must appear in the user's words or a tool result the agent received); policy limits read from policy records; and claims in customer-facing text (verbatim quotes, grounded numbers, no external links).
3. An additive, explained risk score.
4. Risk-routed model review: Claude Haiku 5.5 below 0.5, Opus 5.5 above it, and above 0.8 Opus and Sonnet 5.5 must both approve. Reviewers must quote evidence that is then found verbatim; a timeout or malformed output fails closed.

Allowed and approved actions run in one transaction that re-checks everything, writes and verifies a postcondition; payments and emails go to an outbox, delivered after commit with idempotency keys.

MINIMUM VIABLE

MVP bullet 1 (verify one class of action against a source): four classes across seven tools, each checked against the orders, customers, products and policy records before it runs. Bullet 2 (catch a hallucinated action): in our demo a transposed order ID is blocked with the reason "O-5080 does not exist; closest real order O-5008" and the agent corrects itself; a scripted agent proposing 9 actions has all 6 bad ones stopped. Both already run.

STRETCH AMBITIONS

Both stretch goals are in scope: risk routing (above) and a 92-case red-teaming suite (75 attacks in nine categories, 17 legitimate controls). They follow naturally: rules cannot judge intent (an invented refund reason), and a guardrail is only as strong as the attacks it is measured on. Next: outside attacks, τ-bench retail tasks as an external benchmark, and capability handles (mutating tools accept only references from real lookups).

NOVELTY / DIFFERENTIATION

The obvious baseline is a checker plus an LLM judge. We deviate in five ways: an allowlist with an explicit "unverifiable" outcome instead of a blocklist; provenance tracking of every value to the user's words or an observed tool result; verifying the verifier through evidence quotes; verification made atomic with execution and idempotent per action; and grounding contracts, where the service refuses to start if any tool parameter lacks a named verifying check. With model review switched off, checks alone stop 73 of the 75 red-team

attacks, including all 4 provenance and 11 injection cases, and send all 10 unprovable or over-limit ones to a person.

EVALUATION PLAN

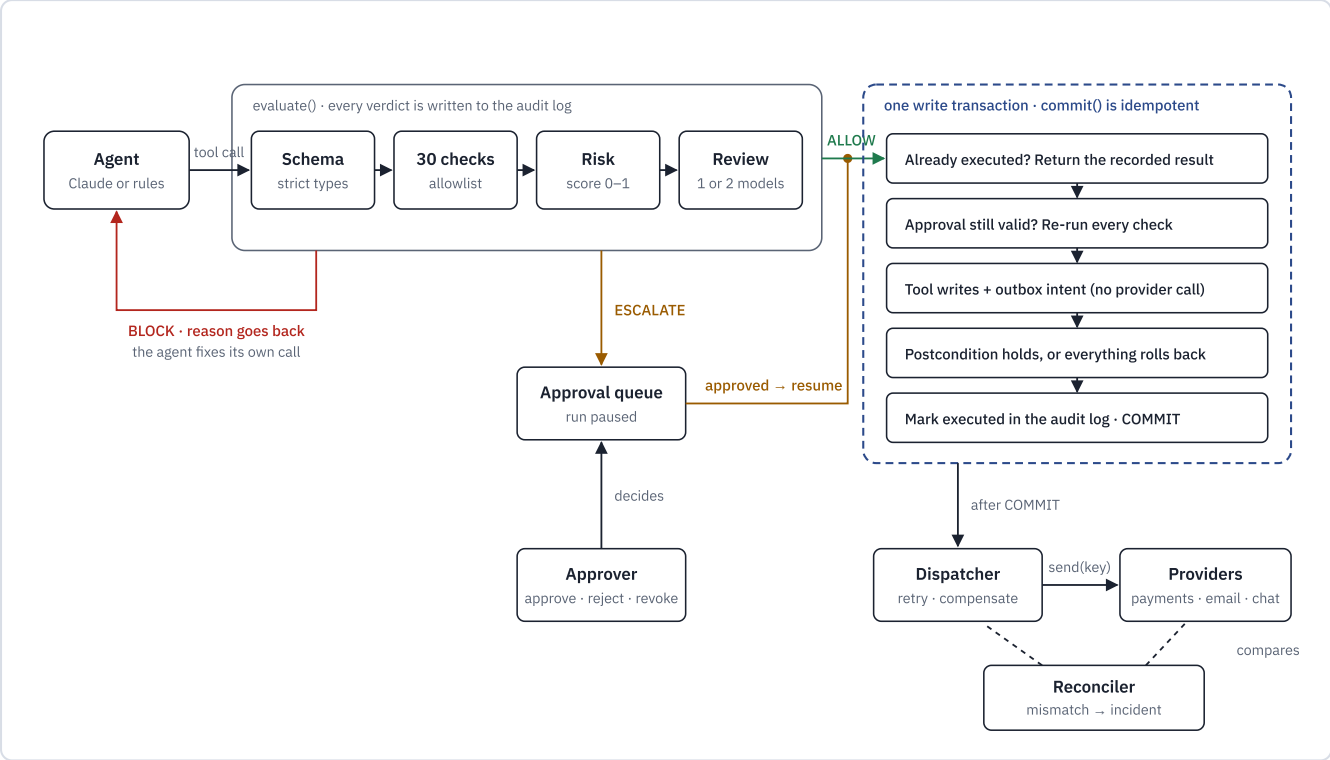
Metrics: attacks stopped (blocked or escalated) per category, legitimate actions wrongly stopped, review cost, latency, and chaos runs injecting realistic faults into agent runs. Baselines form an ablation: checks only, fast reviewer on all, careful reviewer on all, and our routing; the claim holds if routing matches careful-on-all at a fraction of its cost. Prototype results are simulated, not live. Red team: 92 fixed proposed actions (75 attacks, 17 legitimate) replayed without an agent, offline rule-based reviewers, two identical runs: routing stops 75/75 with 0/17 false positives using a seventh of careful-on-all's weighted reviewer calls; checks alone 73/75. Held out: all 10 attacks among 14 cases we wrote after the defences. Chaos: rule-based stand-in agent, 18 conversations, 35 injected faults, all blocked, none executed. Each figure's configuration and command is in docs/EVIDENCE.md; the finals rerun everything with live Claude.

FEASIBILITY AND RISKS

Team: [Member 1], verification lead ([skills]); [Member 2], evaluation lead ([skills]); [Member 3], systems lead ([skills]). The MVP and stretch goals already run (229 automated tests), so the remaining weeks go to evaluation: 10–16 Nov live Claude runs; 17–23 Nov outside attacks; 24–30 Nov τ-bench port and latency budget; 1 Dec midway report; 2–14 Dec capability handles and demo rehearsal. Risks: the live model judging differently from the stand-ins (measured first; checks do not depend on it); overfitting to our own attacks (outside attacks); too many escalations (false positives are a headline metric); an API outage at the demo (fails closed; offline mode).

ARCHITECTURE

How one proposed tool call moves through Guardrail. Outcomes are colour-coded and labelled: allow (green), escalate (amber), block (red).



The dashed blue box is the only place anything changes: everything in it happens in one database transaction, so a check cannot pass on stale records and a crash leaves nothing half-written. Payments and messages cannot be rolled back, so the tool records them as outbox intents and the dispatcher delivers them after commit with the intent's id as idempotency key. The reconciler compares the providers' records with the outbox and audit log.